


BAYLOR ACADEMY FOR TEACHING & LEARNING • ENGINEERING & COMPUTER SCIENCE

Teach for Integrity

Teaching for learning integrity — not just academic integrity — in the age of AI

Megan Alstot, Ph.D. • Academy for Teaching & Learning



A note before we go further.

“You didn't ask for AI to arrive at our educational doorstep. . .

You didn't ask for an unregulated, untested, unreliable, untraceable, and unrelenting technology to upend education... it is clear the genie is out of the bottle... it is not going back in nor is it going to slow down for us to catch up” (Ludwig & Zakrajsek, 2026, p. 3).

Neither did your students.

— *With language from Ludwig & Zakrajsek (2026), The Science of Learning Meets AI*

Forty-five minutes, one big topic, one thing you can use in week one



Disrupting Routines: Cognition, Fluency, Literacy

What is actually under threat, and the frame this session uses to talk about it.



Weeks one and two

The highest-opportunity, lowest-tech integrity builders.



Why HELP is a four-letter word

What happens when a student asks AI for help — and who ends up holding the decision.



Mentor, not police

Four small teaching approaches to building integrity.



Take home a toolkit

Your gift: a page of ready-to-use tools (and maybe time to try them).



[malstot.github.io/
gifts/ecs-2026](https://malstot.github.io/gifts/ecs-2026)

Scan the QR any time — the toolkit is open from now on, not just at the end.

Where are you with AI — in your course?

Show of Hands

Hold up fingers — one through five.

1

2

3

4

5



Hesitant — I'd rather AI stayed out of my course

Confident — AI already has a seat in my course

Wherever your hand landed, we can probably agree on this much: learning — students' cognitive development — is the point. AI threatens that thinking when it's used thoughtlessly.

What's the kernel — and what's at stake?

Every course exists to build some kind of thinking necessary for success in your discipline — either as a scholar or a practitioner.

You are best positioned to decipher what that thinking is. That is the Kernel of Cognition.

1 Only the discipline expert can identify the Kernel.


No campus policy can decide which cognitive work is essential in your course — the derivation, the debugging, the design judgment, the proof. You can. And only you can determine the potential of AI to support or subvert that learning.

2 Cognitive offloading isn't automatically bad.


Sometimes it frees capacity for harder thinking. Sometimes it removes the very struggle the course was built around. The difference is the role AI is invited to play.

3 The routine is the risk.

Students build AI routines shaped by their social context. Once a routine goes tacit, almost no real thinking survives inside it.



THE FULCRUM



The kernel of cognition

The thinking, growing, and learning a degree exists to build. If an AI use short-circuits it, the kernel is surrendered — no matter how good the output looks.

A Critical Core for AI in Education

Each one is a lens for thinking and AI.


Cognition is the fulcrum point between enhancing and surrendering learning.

Fluency without attention to **cognition** produces students who over-rely and short-circuit learning.

Using AI without fluency and literacy produces users who can produce with AI but do not know how to regulate it.

Take any one away, and the other two can't carry the weight.


At the center: all humans in the learning loop. Students • Faculty • Leadership



My focus: keep the human in the learning loop.

HITL² — Humans in the Learning Loop

You're a human in a learning loop, teaching humans in your learning loops — and every one of us is figuring out how to learn in an AI-integrated loop. Together.



You



Your Students



All of Us

Why teach for fluency?

Fluency doesn't presume AI should be used — full stop. Fluency builds ethical decision making.

IF

Should AI be part of this at all? “No, not here” is a fluent answer or “Yes it can ...” are both fluent answers..

WHEN

At which stage? Before the thinking starts, in the middle of it, or only after the student has produced something.

HOW

In what role, and inside what rails? Coach, critic, and ghostwriter are not the same invitation.


WHY

What does it serve — and can the student say what they learned that they wouldn't have otherwise?

We CAN'T prepare them for the AI-integrated world they'll navigate in 10 years (10 weeks?)

We CAN prepare them to THINK about how they integrate AI in it.

The 4Ds of AI Fluency — a frame for thinking



Delegation

Setting goals and deciding whether, when, and how to engage AI at all.

Description

Describing your goals clearly enough to prompt useful AI behavior — guardrails included.

Discernment


Accurately assessing the usefulness of AI outputs and behaviors.

Diligence

Taking responsibility for what we do with AI — and how we do it. Responsibility rests with the human.

The routine is the risk

Students, AI, faculty and social context are one co-constituted field — you cannot stand outside it



PART ONE • LOW TECH, HIGH IMPACT

The first two weeks are your highest-opportunity weeks

This is when you build relationships and establish shared values.

Eleven reasons students compromise their learning integrity

IN YOUR GIFT



Self-efficacy — Doubts they can succeed, so reaches for a shortcut.



Expectations — One instructor's “AI feedback” is another's “plagiarizing AI.”



Values — The task feels disconnected from what they came for.



Motivation — Playing for the grade, not the learning.



Procrastination — Out of time; AI is the only door still open.



Cost-benefit — Weighs it: high upside, low chance of getting caught.



Moral disengagement — Rationalizes it — “everyone does it.”



Contract cheating — Hands the work to someone — or something — else.



Peer effects — Takes cues from what the class treats as normal.



Narcissism — Acts as if the rules are for other people.



Instructor effects — Doesn't feel the instructor knows or cares.

Non-technical, low effort, biggest net. “Talk is cheap” — discursive moves are the lower long-term lever (Corbin et al., 2025) — but they're the easiest on-ramp and a means for building trust through relationships.



Learner survey

Meet students week one: names, goals, how they already use AI. (Weeks 0-2)



Integrity conversation

A real talk about integrity when the first higher-stakes task is set. (Not day 1!)



Integrity nudges

Reminders at announcement and just before the due date.



Grace period

One no-questions 48-hour extension (cap the uses). — Clobes, CSUCI



Withdraw & take the 0

24 hours to pull a submission, no questions — restores moral agency. — Bowen & Watson, 2025

Builds against the widest net — 8 of the eleven (self-efficacy, expectations, motivation, procrastination, cost-benefit, moral disengagement, peer & instructor effects).

Assume good intentions — and limited AI literacy.

“[Generative AI] represents a fundamental shift in the relationship between students and technology in their learning. Alongside the new learning practices that students are engaging in, new teaching practices and policies are needed. Adding more content and/or focusing on routine administrative processes to stop cheating are well short of what will be required to adapt.”

— Lodge, Thompson, & Corrin (2023a), p. 7

So — is “cheating” with AI a character problem, or a fluency problem?

PART TWO • THE NEW FOUR-LETTER WORD

HELP

What happens when a student asks AI for help?

Watch what four words can do.

The ask

Help me improve my learning objective.

Standards

California ELA/Literacy Content Standard(s) (CA CCSS)

W.4.3.A

Orient the reader by establishing a situation and introducing a narrator and/or characters; organize an event sequence that unfolds naturally.

W.4.3.B

Use dialogue and description to develop experiences and events or show the responses of characters to situations.

Learning Goals/Objectives

At the end of this lesson the Students will be able to (TSWBAT):

[Students will be able to explain characteristics and emotions of characters in the story by looking back upon text in the story "Island of the Blue Dolphins."]

Assume good intentions and limited How-to-Description Skills.

It answered & offered

Your objective is clear and student-friendly, but there's a mismatch worth fixing first: **W.4.3.A** and **W.4.3.B** are narrative *writing* standards, and your objective asks students to *analyze* a text. "Explain characteristics and emotions of characters by looking back at the text" is really RL.4.1 (refer to details and examples when explaining) and RL.4.3 (describe a character in depth, drawing on specific details).

So you have two clean ways forward.

Option 1 — Keep the objective, swap the standards to reading:

TSWBAT describe two of Karana's character traits and the emotions she experiences during [specific chapter/event], citing at least three specific details from *Island of the Blue Dolphins* to support each inference, in a written paragraph.

Option 2 — Keep the W.4.3 standards and make the objective a writing task:

TSWBAT write a narrative scene from Karana's point of view that establishes the situation and introduces her as narrator (W.4.3.A), using at least two lines of dialogue and three sensory details to show how she responds to [the event] (W.4.3.B).

Option 3 (my recommendation) — Bridge them, since this is the lesson you seem to be teaching:

TSWBAT analyze Karana's traits and emotions by citing at least three specific details from *Island of the Blue Dolphins* (RL.4.3), then use those details to write a short narrative scene from her perspective that develops her response to the event through dialogue and description (W.4.3.A-B).

What changed and why:

- **Named the character.** "Characters in the story" is broad; Karana gives students a target and makes the lesson assessable.
- **"Characteristics" → "traits."** Matches the academic vocabulary in the standard and in most 4th-grade curricula.
- **"Looking back upon text" → "citing specific details."** Same idea, but it's the standard's language and it tells students what the behavior actually looks like.
- **Added a product and a degree.** "Explain" doesn't tell you what you'll collect. "Write a paragraph with three details" is something you can hold in your hand at the end of the period and know whether they got it.
- **Bumped the rigor.** "Explain" sits at DOK 2. "Analyze... then use those details to write" moves into DOK 3, which is where the W.4.3 standards live.

One thing to check before you commit: is the writing piece something students will actually finish in this lesson, or is it the next day? If it's a two-day arc, I'd split this into two objectives — reading Day 1, writing Day 2 — rather than stretching one objective across both. Want me to draft the success criteria or a quick rubric to go with whichever version you pick?

"Want me to draft the success criteria or a quick rubric to go with whichever version you pick?"

The ask with human-regulation

Context: I am a student preparing to be an elementary school teacher in CA. I need to work on writing measurable learning objectives for my lesson plans

Assignment: I am developing a lesson plan for 4th graders reading Island of the Blue Dolphin. The first part of the lesson plan must have measurable learning objectives that align with the ELA standard I've selected.

Role: You are an expert teacher. You are going to check my learning objectives and if needed, help me improve them so that they align with the ELA standards, are measurable, and appropriate for 4th-grade students (typically ages 9 to 10 years old)

Example: Do not rewrite my learning objectives for me. Instead, identify which objectives need work. Then act as a coach to help me improve them on my own. You can ask me questions, tell me which part of my objective needs revising, and point me to resources that can help me. You can tell me when I have provided a measurable learning objective. You can help me think of ideas for gathering evidence that students have met the objective. But you should not make any decisions for me. I am the only one who can decide what the final objectives are. You can evaluate whether they are measurable and align with the ELA standard.

Cognitive Load! Give me one or two suggestions at a time. Don't overwhelm me.

How to start: When you are ready, ask me to share my ELA standard and my learning objectives.

The ask with human-regulation

Context: I am a student preparing to be an elementary school teacher in CA. I need to work on writing measurable learning objectives for my lesson plans
Assignment: I am developing a lesson plan for 4th graders reading Island of the Blue Dolphin. The first part of the lesson plan must have measurable learning objectives that align with the ELA standard I've selected.
Role: You are an expert teacher. You are going to check my learning objectives and if needed, help me improve them so that they align with the ELA standards, are measurable, and appropriate for 4th-grade students (typically ages 9 to 10 years old)
Example: Do not rewrite my learning objectives for me. Instead, identify which objectives need work. Then act as a coach to help me improve them on my own. You

Show more

I'm ready when you are.

Please share two things:

1. **The CA ELA standard** you've selected (the full text, or the code plus your paraphrase — whichever you have).
2. **Your learning objectives** as you've currently written them.

Once I see both, I'll tell you which ones are already measurable and well-aligned, and which need work — one or two at a time so we can take it in manageable pieces.

Standards

California ELA/Literacy Content Standard(s) (CA CCSS)

W.4.3.A

Orient the reader by establishing a situation and introducing a narrator and/or characters; organize an event sequence that unfolds naturally.

W.4.3.B

Use dialogue and description to develop experiences and events or show the responses of characters to situations.

Learning Goals/Objectives

At the end of this lesson the Students will be able to (TSWBAT):

[Students will be able to explain characteristics and emotions of characters in the story by looking back upon text in the story "Island of the Blue Dolphins."]

Thanks — this gives me plenty to work with. Let's start with the biggest thing, and we'll get to measurability after.

The alignment issue

Look closely at what your standards ask students to *do*:

- W.4.3.A — "**Orient the reader...** introducing a narrator and/or characters; **organize an event sequence**"
- W.4.3.B — "**Use dialogue and description** to develop experiences... or show the responses of characters"

Both are W standards — Writing. The verbs describe students *producing* a narrative.

It handed the decision back

Both are W standards — Writing. The verbs describe students *producing* a narrative.

Now look at your objective: "**explain** characteristics and emotions of characters... by **looking back upon text**." That's students *analyzing* a text someone else wrote. That's reading comprehension work — closer to RL.4.1 (referring to details and examples in a text) or RL.4.3 (describing a character in depth, drawing on specific details).

So right now the objective and the standard are doing two different jobs. Neither one is wrong — they just don't match yet.

This is your call, and there are two honest paths:

1. **Keep W.4.3.A/B** and revise the objective so students *write* something — using dialogue and description to show how a character responds to a situation.
2. **Keep your objective** and swap to the RL standard that actually describes character analysis.

There's also a third option teachers use often: the reading analysis becomes the *scaffold*, and the writing is the objective — students study how O'Dell shows Karana's fear or determination, then use that technique in their own narrative writing.

Which of those feels closest to what you actually want kids doing at the end of this lesson? Tell me that, and we'll work on the wording from there.

Both versions found the same problem. Only one left the decision with the student.

An example of co-regulation (Lodge et al., 2023).

"I explicitly tell it, do not write this for me."

"I'll typically go to ChatGPT and specifically ask for it not to generate — which is what it thinks I want, which is to just write the essay for me. I explicitly tell it, do not write this for me. I need to do this on my own, but I want help with understanding what I should write about" (Ian, Y1 • Computer Science).

He wrote a guardrail. How do we build co-regulation routines in students?

— Alstot (2026), dissertation interview data. Co-regulation is a GenAI literacy — and it is teachable.

PART THREE • THE PRACTICAL PART

Mentor, not police

Four moves inside your own sphere of influence. Pick ONE to try this term — the rest live in your gift.

- 1 Don't → Do pairs
- 2 Stress test
- 3 Write your rail
- 4 RADAR — student rails

Cheating has always existed. AI changed how — not why.

Policing asks:

“How do I catch them after?”

- Surveillance, suspicion, detection tools
- Exhausting — and error-prone with AI
- Erodes the relationship that prevents cheating

Teaching for integrity asks:

“What invites the shortcut — and how do I design it out?”

- Conversation, early and often
- Clarity about the line, and the why behind it
- Design that makes thinking visible

The real struggle isn't “academic integrity.” It's learning integrity. *We want students to learn — so the charge is to teach for integrity: learning and academic.*

Bertram Gallant & Rettinger (2025)

Every Don't needs a Do

A prohibition on its own just stalls the conversation. Pair it with the move you want instead.

Don't

Generate or rewrite the section for the student.

Give the answer.

Offer to fix it for the student.

Do

Ask questions about...

List gaps or inconsistencies in...

Identify sections to strengthen or shorten...

Write one pair for the assignment you're teaching soonest. A prohibition alone just stalls the conversation.

Stress test the assignment



How does AI do on your assignment?

Paste your assignment into the Assignment Integrity Stress Test: seven students, each on a bad day, built from the eleven reasons.

It diagnoses; it does not fix.

- You get a misuse-pressure report and a question back: which weakness do you want to work on, and how?
- Where is AI strong and students are weak? Can you integrate AI here?

It shows you two things at once: what guardrails this assignment needs, and where you need to consider redesigning the assignment

Builds against:

all eleven — it diagnoses

Copied — paste it into a new Copilot or Claude chat.

▼ Preview the full prompt

You are running an ASSIGNMENT INTEGRITY STRESS TEST. Do not begin until I have pasted my assignment. First ask me for: (1) the assignment text as students see it; (2) the kernel of cognition – "the thinking this assignment exists to build is ____"; (3) its AI literacy level: Level 1 AI-Free Assignment (deliberate, task-specific – not a course-wide ban), Level 2 Limited (narrow tasks only), Level 3 Collaborative (I draft, AI critiques), Level 4 Integrated (AI as professional tool), or Level 5 Innovation (students design novel uses with self-set rails); (4) discipline, course level, and modality; (5) anything I know about my students (aggregated learner-survey patterns – confidence, which AI tools they use; never identifiable data).

Then stress-test the assignment through 7 student personas, each on a bad day, built from the eleven research-backed reasons students cheat (Bertram Gallant & Rettinger, 2025):

1. THE DOUBTER (self-efficacy): doesn't believe they can do this; AI feels like survival. Looks for practice and feedback before stakes rise.
2. THE LINE-GUESSER (expectations): genuinely unsure what counts as cheating here; reads silence or vague policy as permission.
3. THE EFFICIENCY RATIONALIZER (values + moral disengagement): a working professional who sees busywork – no visible connection to what they're here to develop as a scholar or practitioner; "everyone does it" ready to go.
4. THE 11 PM PROCRASTINATOR (procrastination + cost-benefit): deadline tonight; a full work week and family obligations ate the runway; calculates upside high, risk low, relationship thin.
5. THE GHOST-WRITER SHOPPER (contract cheating): will submit AI output wholesale if only the final product is ever seen; authorship undefined.
6. THE NORM READER (peer effects + instructor effects): takes cues from classmates and from whether the instructor knows them and reads their work.
7. THE GRADE MAXIMIZER (motivation + narcissism): plays for the grade; suspects rules are for others; exploits any product-only design.

Write the rail your stress test asked for



You don't have to invent rails from nothing. Run the stress test first: what it exposes is what your rail has to cover.

Write it:

Start from the guardrail starter handout in your gift — example rails you adapt, not a template you fill in. Each one pairs the rail with a student-facing why.

Place it:

In the assignment directions

Discuss it:

Orient students to it early, and show them how it works

Builds against:

Expectations

Moral disengagement

Guardrail starter handout — example rails for you and for your students.

The Master Guardrail Prompt

Students should paste this at the start of any AI interaction to establish boundaries:

Master Guardrail — Students Should Use This First

Act as a learning coach. Do not write, solve, summarize, or revise my work for me. Do not generate final answers. Instead, help me think by asking questions, identifying issues, giving feedback, and suggesting next steps.

This one prompt reframes the entire interaction before it begins. Encourage students to make it a habit — paste first, then ask.

Managing Pacing & Cognitive Load

Students can add this to any prompt to prevent AI from overwhelming them:

Pacing Guardrail

Ask me one question at a time. Make one suggestion at a time. Wait for my response before continuing.

Limiting AI to Student-Provided Materials

When students want help with their own notes, readings, or drafts — not outside content:

Materials-Only Guardrail

Use only the materials I provide. Do not draw from outside sources. If information is missing, tell me what I need to supply — do not fill the gaps yourself.

A student-facing prompting scaffold that keeps students working within the continuum level you set. It's how they operationalize Delegation and Description.

R Your role

"I'm a student working on [assignment]..."

A AI's role

"Act as a peer reviewer / study coach / research librarian."

D Describe the task

What you're working on — and the sources or notes to use.

A Action to delegate

"Flag gaps / generate questions / check my logic — [the specific ask]."

R Rails — the guardrail

"Don't write it for me. Flag, don't fix — I'll decide. Ask before you suggest."

The last R is where integrity lives — the rail students set before they ask.


Builds against:

Expectations

Moral disengagement

Where should this assignment sit? A continuum, not a scale

IN YOUR GIFT



Every point is a real choice

- Each level implies a role for AI and the literacy students need to work fluently at that level.
- How can the assignment design develop these literacies?

Course v. Assignment

- You can have a "No-AI" Assignment. You cannot have a "No-AI" course
- Even "No-AI" assignments can teach AI literacy

Supports MLO A4.1–A4.2 (the five levels, how they build, and the AI literacy each develops). AI-literacy focus adapted from the IWU N&G AI Literacy and Integration Scale (Bleistein & Snyder, EDUCAUSE Review, 2026); continuum adapted from the IWU 5-Level AI Literacy Continuum (Indiana Wesleyan University).

Your “gift” — A toolkit to take with you

OPEN IT ANY TIME

1 Learner Information Survey

Meet students in week one: names, goals, how they already use AI.

2 The AI Literacy Continuum

Place an assignment. Every level is a real choice when it's intentional.

3 Assignment Integrity Stress Test

Paste an assignment, meet your seven bad days. Four minutes, no prep.

4 Guardrail starter handout

Example rails to adapt — for your assignment and for your students.

5 Draft Your Disclosure Stance

Four questions, and you leave with syllabus language in your own voice.

6 The Eleven Reasons

The full mentoring matrix behind the slide — the research and the response.



[malstot.github.io/
gifts/ecs-2026](https://malstot.github.io/gifts/ecs-2026)

The page is yours to keep, share with your department, and come back to all term.

Keep the human in the learning loop.

These forty-five minutes were the start of a conversation. Leave with one thing — and know where to find me for the rest.

Turn to someone near you: the ONE move or strategy you'll try this term.

That's what I'm here for. Consultations · small-group sessions · workshops built for your unit or team.

megan_alstot@baylor.edu



References, attribution & license

- Alstot, M. (2026). Generative artificial intelligence and the metaliterate learner: A mixed methods study of AI-literacy and US undergraduate students' interaction within an AI-integrated knowledge economy (Publication No. 32582165) [Doctoral dissertation, Pepperdine University]. ProQuest Dissertations and Theses Global.
- Alstot, M. N., Goldberg, D. M., Sobo, E. J., & Hauze, S. W. (in press). A sociomaterial view of AI diffusion in higher education: How institutional lag thwarts responsible AI use among students as early adopters. *Communications of the Association for Information Systems*.
- Bertram Gallant, T., & Rettinger, D. (2025). *The opposite of cheating: Teaching for integrity in the age of AI*. University of Oklahoma Press.
- Bleistein, T., & Snyder, T. (2026, March 31). Advancing AI literacy: A faculty course refresh institute at Indiana Wesleyan University. *EDUCAUSE Review*.
- Corbin, T., Dawson, P., & Liu, D. (2025). Talk is cheap: Why structural assessment changes are needed for a time of GenAI. *Assessment & Evaluation in Higher Education*, 50(7), 1087–1097.
- Dakan, R., Feller, J., & Anthropic. (2025). *AI fluency: Framework & foundations*. CC BY-NC-SA 4.0.
- Fawns, T. (2022). *An entangled pedagogy: Looking beyond the pedagogy–technology dichotomy*. Postdigital Science and Education.
- Lodge, J. M., Thompson, K., & Corrin, L. (2023a). Mapping out a research agenda for generative artificial intelligence in tertiary education. *Australasian Journal of Educational Technology*, 39(1), 1–8.
- Ludwig, D. L., & Zakrajsek, T. D. (2026). *The science of learning meets AI*. Routledge.
- Major, C. H. (2026, May). Critical competencies in the age of AI [Keynote]. Lilly Conference, Austin, TX.
- Perkins, M., Furze, L., Roe, J., & MacVaugh, J. (2024). The Artificial Intelligence Assessment Scale (AIAS). *Journal of University Teaching and Learning Practice*, 21(6).
- Winkel, D. (2026). Stop reviewing your lesson plans alone. *How to Teach with AI*.

USE, ADAPT, SHARE — WITH CARE

Original workshop materials © 2026 Megan Alstot, Academy for Teaching & Learning, Baylor University — licensed CC BY-NC-SA 4.0 (adapt for non-commercial educational use, with attribution, shared alike).

NOT covered: the cited scholarship above (cite the sources); the Ludwig & Zakrajsek excerpt (quoted, not re-licensed); the 4D Framework (Anthropic, CC BY-NC-SA 4.0); the continuum (adapted from IWU / Bleistein & Snyder 2026, informed by AIAS / Perkins et al. 2024); Winkel's review pattern; Baylor and HITL²™ names and marks.